

Risk-Aware Offline Reinforcement Learning with Calibrated Off-Policy Evaluation for Motor-Imagery BCI Decoding

Anonymous Submission

Paper ID #XXX

Author identity withheld for double-blind review

Abstract—Motor-imagery brain-computer interfaces commit to a class label at a fixed window after the cue, paying the full wait cost on every trial and offering no principled way to abstain when the signal is ambiguous. We reframe within-trial decoding as a sequential decision a learned agent makes at each 0.25-s timestep: commit a class, defer for more evidence, request recalibration, or abstain. The agent is a Conservative Q-Learning controller trained offline on logged trials under a CVaR risk-sensitive critic and a CMDP commit-error constraint that directly bounds the rate of unsafe decisions. We validate it with a calibrated fitted-Q off-policy evaluator that selects policies without redeployment on humans. On a 12-policy panel, the estimator tracks on-policy Monte-Carlo ground truth at Pearson $r = 0.860$ with rank ordering preserved, sufficient to pick the best of a candidate set without an online trial — the *first* calibrated OPE result published for motor-imagery BCI. On BCI-IV-2a 4-class with an EEG-Conformer encoder, the selective decoder reaches $87.6 \pm 2.9\%$ commit accuracy at 37.0 bits/min ITR on the canonical 3-subject panel, exceeding the mandatory EEGNet baseline (0.815 / 20.7 b/m) on every axis; extended to all 9 subjects it reaches $74.6 \pm 11.9\%$ at 24.9 b/m. Per-subject results cluster bimodally on commit accuracy: four strong subjects match the headline (0.84–0.89), four hard subjects cap at 0.60–0.69 regardless of agent, and their lower encoder validation accuracy (≤ 0.66) tracks the floor — locating future gains in encoder quality, not policy machinery. The CMDP commit-error budget ($\epsilon = 0.10$) is satisfied on every reported configuration, and a cross-subject CVaR α sweep exposes the risk knob as a designer-tunable safety-throughput trade-off (commit rate $0.16 \rightarrow 0.29$, wrong-rate $0.022 \rightarrow 0.040$ across $\alpha \in [0.10, 1.0]$) — a clean operating-point dial clinicians and regulators can specify in advance. The framework transfers across channel counts. Leave-one-subject-out classification on 62-channel Lee2019 delivers the *first* published LOSO benchmark for offline-RL BCI (10/10 subjects beat random, paired Wilcoxon $p = 0.001$). Against an EEGNet-fixed mandatory baseline on the same splits, commit accuracy is statistically comparable (0.696 vs 0.716, $p = 0.43$) while wrong-commit rate trends 25% lower (0.214 vs 0.284, $p = 0.07$, marginal) — a safety advantage rather than an accuracy SOTA. To our knowledge this is the first offline-RL BCI decoder with DEFER as a *learned* action; together with the calibrated OPE methodology, it provides the missing infrastructure for closed-loop clinical BCI deployments where wrong-decision cost is bounded by design and policy choice is auditable offline. Code released under MIT.

Index Terms—brain-computer interface, motor imagery, offline reinforcement learning, off-policy evaluation, fitted Q-evaluation, Conservative Q-Learning, EEG decoding, sequential decision making.

I. INTRODUCTION

A motor-imagery (MI) brain-computer interface (BCI) must decide *what* the user intends and *when* to commit that decision. Standard pipelines (FBCSP+LDA, EEGNet, EEG-Conformer) decode at a fixed window—typically 0–4 s post-cue—trading accuracy against information transfer rate (ITR), the throughput metric that matters for end-users. Dynamic-stopping classifiers [4], [27] and SPRT variants [17] introduce evidence-based early termination, but these are hand-tuned thresholds on confidence statistics, not learned policies.

Recent work [3], [7], [22] frames BCI decoding as sequential decision-making, but none use offline reinforcement learning, none expose an explicit DEFER action as part of the learned policy, none address motor imagery at scale, and—critically—none employ a calibrated off-policy evaluator. Existing evaluations simulate a single classifier under a fixed protocol, which is on-policy simulation in a deterministic environment, not OPE in any rigorous sense. Meanwhile, offline RL with calibrated OPE has matured in adjacent clinical domains (sepsis, ventilator weaning, anesthesia depth control) [12], [14], [23], [28], but this machinery has not been ported to BCI.

This paper fills that methodological gap. We frame within-trial MI EEG decoding as an offline RL problem with a clinically meaningful action space $\{\text{COMMIT}_k, \text{DEFER}, \text{RECAL}, \text{ABSTAIN}\}$, train a Conservative Q-Learning (CQL) agent on logged trajectories from two behavior policies, and validate the resulting policies with a calibrated fitted-Q evaluator against on-policy Monte-Carlo ground truth.

Contributions: (i) **Calibrated OPE for BCI:** per-target-policy FQE attains Pearson $r = 0.860$ on a 12-policy panel versus on-policy Monte-Carlo — the first published calibrated OPE result for BCI motor imagery. (ii) **Selective offline-RL decoder:** a CVaR critic combined with a CMDP commit-error constraint reaches commit accuracy $87.6 \pm 2.9\%$ at 37.0 bits/min ITR on the canonical 3-subject BCI-IV-2a panel (EEG-Conformer encoder) and $74.6 \pm 11.9\%$ at 24.9 b/m on the full 9-subject panel (Table III) — the first offline-RL BCI decoder with DEFER as a *learned* action (distinct from post-hoc thresholding [8]). (iii) **Cross-subject LOSO on three datasets:** 62-channel Lee2019 ($N = 10$) yields

commit accuracy 0.696 ± 0.104 vs random 0.481, 10/10 wins, paired Wilcoxon $p = 0.001$ — the first published Lee2019 LOSO benchmark for offline-RL BCI, completing a monotonic channel-count to LOSO-significance trend ($p = 0.10/0.15/0.001$ on $\{3, 22, 62\}$ -channel datasets). Against an EEGNet-fixed mandatory-commit baseline on the same splits, raw commit accuracy is statistically comparable ($p = 0.43$); the agent’s advantage is a -25% relative reduction in wrong-commit rate (0.214 vs 0.284, $p = 0.07$), positioning the contribution as a safety-oriented selective decoder rather than an accuracy SOTA. (iv) **Foundation-model RL fine-tuning:** first end-to-end offline-RL fine-tuning of the LaBraM [11] EEG foundation model on a BCI task; ties EEGNet in aggregate (paired Wilcoxon $p = 0.79$) with a reproducible subject-specific encoder-rescue effect (+16.8pp on BCI-IV-2b, +15.9pp on BCI-IV-2a for the same held-out individual). (v) **Open release** of the entire pipeline (data, MDP simulator, agent, OPE estimators, LaBraM integration) under the MIT license.

Topic match: The work falls in two named topical areas of the venue: *Adaptive Control* and *Neural Signal Processing*. The methodological contribution—turning a within-trial decoding decision into a learned, safety-constrained sequential decision problem under a calibrated OPE estimator—is a direct application of intelligent control to neural-signal-processing pipelines.

II. RELATED WORK

Dynamic stopping in BCI: Verschore et al. [27] introduced classifier-confidence thresholds for early termination on a P300 speller; Bianchi et al. [4] proposed a refined early-stopping decision rule in the same paradigm. Liu et al. [17] adapted a sequential probability ratio test (BTSPRT) to motor-imagery BCI. All are hand-tuned thresholds on a confidence statistic: there is no learned policy and no recalibration or abstention action.

Sequential decision frameworks for BCI: MARKOV-TYPE [22] casts a P300 RSVP speller as a POMDP solved with on-policy REINFORCE. EEG_RL-NET [3] trains a Dueling DQN classification head on a GNN backbone, online and without a DEFER action. ErrP-driven online RL [7] adapts a decoder during a paced game using error-related potentials as the reward signal. None of these are *offline* RL; none use a DEFER, RECAL, or ABSTAIN action; none ship a calibrated OPE estimator.

Offline RL and OPE methodology: Conservative Q-Learning [12] and risk-sensitive distributional variants [6] are now standard; per-decision importance sampling [18], doubly-robust estimators [10], and fitted-Q evaluation [14] are the workhorse OPE estimators. Our contribution is porting this stack to BCI and validating it against on-policy Monte-Carlo ground truth on a deterministic per-trial environment.

EEG foundation models: LaBraM [11] (ICLR’24 spotlight, $\sim 2,500$ hours of EEG pretraining), CBraMod, and EEGPT release pretrained encoders; every published downstream usage attaches a supervised classification head. Our

cross-subject experiments (§IV-J) are the first end-to-end RL fine-tuning of an EEG foundation model in BCI.

III. METHODS

A. Problem Formulation

A motor-imagery (MI) trial is a multivariate EEG segment $X \in \mathbb{R}^{C \times N}$ with C channels, N samples, and class label $y \in \{1, \dots, K\}$. We view single-trial decoding as a sequential decision problem: at each timestep, the system must choose to commit a class label, defer for more evidence, request a recalibration trial, or abstain. The augmented action set is $\mathcal{A} = \{\text{COMMIT}_1, \dots, \text{COMMIT}_K\} \cup \{\text{DEFER}, \text{RECAL}, \text{ABSTAIN}\}$.

We discretise the trial into windows of length $L = 1.0$ s with stride $\Delta = 0.25$ s, yielding $T = \lceil (N/f_s - L)/\Delta \rceil + 1$ window indices per trial ($T = 13$ for a 4 s trial at 250 Hz). The per-trial Markov decision process (MDP) is defined as follows.

State: s_t is the concatenation of (i) a window encoding $\phi(X_t) \in \mathbb{R}^D$ from a pretrained encoder ϕ ; (ii) the running mean and variance of $\phi(X_{1:t})$; (iii) elapsed time $t/(T-1)$; (iv) a recalibration flag $r_t \in \{0, 1\}$; and (v) a D_s -dim subject embedding e_{subj} . Total $\dim(s_t) = 3D + 2 + D_s$.

Transitions: Deterministic given the trial recording. DEFER advances $t \mapsto t + 1$. COMMIT and ABSTAIN terminate. RECAL pays a fixed time cost τ_r (advancing t by $\lceil \tau_r/\Delta \rceil$ steps) in exchange for replacing e_{subj} with a per-subject embedding refit from a held-out calibration buffer; $r_t \mapsto 1$ and RECAL becomes unavailable.

Reward: Per-second time cost $\lambda = 0.05/\text{s}$; correct commit $+R^+ - \lambda t\Delta$ ($R^+ = 1$); wrong commit $-R^- - \lambda t\Delta$ ($R^- = 5$); DEFER $-\lambda\Delta$; RECAL $-R_{\text{recal}} = -0.3$; ABSTAIN $-R_a - \lambda t\Delta$ ($R_a = 0.5$). Constraint $R^- > R_a > 0$ enforces “abstain better than wrong commit”.

B. Encoder and Subject Embedding

We use the EEGNet backbone [13] ($F_1 = 8$, $D = 2$, kernel lengths 64 and 16, dropout 0.25; $\sim 1.5\text{K}$ parameters), supervised- pretrained for 20 epochs on training-subjects’ window-label pairs with a trial-disjoint held-out validation split (10%) for early stopping. The encoder is frozen for the agent.

LaBraM foundation-model integration. We additionally integrate the pretrained LaBraM-base checkpoint [11] (5.8M parameters, pretrained on $\sim 2,500$ hours of EEG) via brain-decode 0.9. The pretrained checkpoint was trained with a 64-channel position- embedding table while braindecode 0.9 defaults to 128; we surgically copy the first 65 position-embedding rows (CLS + 64 channels) and the first 2 temporal-embedding rows (CLS + first patch) from the pretrained checkpoint into our model and accept random init for the remaining channel positions. Input windows (1 s @ 250 Hz) are resampled on-the-fly to 200 Hz to match the patch-size of the pretrained model. We test LaBraM in three modes: (i) frozen body + linear head (LABRAM-FROZEN); (ii) end-to-end supervised fine-tune of the full $\sim 5.8\text{M}$ -parameter model on training-trial windows for 15 epochs of AdamW at

$\text{lr} = 10^{-4}$ (LABRAM-FINETUNED); after which the encoder is frozen and used by the RL agent.

The subject embedding is a fixed Gaussian random projection of dimension $D_s = 16$ applied to the mean encoder feature aggregated across a per-subject set of trials. For LOSO experiments the default and post-recalibration embeddings are pooled across the 8 training subjects, ensuring no held-out-subject information leaks into the embedding.

C. Offline RL Training

We train Conservative Q-Learning [12]. The loss has three terms:

- 1) Bellman residual (Huber): $\mathcal{L}_{\text{TD}} = \mathbb{E}_{(s,a,r,s') \sim \mathcal{D}} [\rho_\delta(Q_\theta(s,a) - y_{\text{tgt}})]$ with $y_{\text{tgt}} = r + \gamma(1-d) \max_{a'} Q_{\bar{\theta}}(s', a')$ and Polyak-averaged target network $\bar{\theta}$.
- 2) Conservative penalty: $\mathcal{L}_{\text{CQL}} = \alpha_{\text{CQL}} (\log \text{sumexp}_a Q_\theta(s,a) - Q_\theta(s, a_{\text{data}}))$.
- 3) Optional CMDP safety penalty: with Lagrange multiplier $\beta = e^{\log \beta}$ and tolerance ϵ , $\mathcal{L}_{\text{CMDP}} = \beta \cdot \mathbb{E}_{(s,a) \sim \mathcal{D}} [\mathbb{I}\{a \text{ wrong commit}\} Q_\theta(s,a)]$ pulls down the value of state-action pairs whose data outcome was a wrong commit. We use $\epsilon = 0.10$ throughout.

The risk-sensitive variant (CVaR-CQL) replaces the Q-network with a quantile-Huber regression à la QR-DQN [6] and uses CVaR $_\alpha$ in policy evaluation while keeping the mean for the Bellman target. The implementation passes a synthetic-bandit unit test (quantiles of $\mathcal{N}(2,1)$ recovered to mean abs error 0.17). Both scalar and CVaR variants, with and without the CMDP constraint, are exercised in the within-subject ablation (§IV-G) and on the canonical BCI-IV-2a 4-class panel (§IV-H).

Optimiser. AdamW, $\text{lr}_{\text{critic}} = 3 \times 10^{-4}$, weight decay 10^{-4} , gradient clip 1.0, batch 256, Polyak $\tau = 5 \times 10^{-3}$, $\gamma = 0.99$. We run 20K gradient steps within-subject and 10K for LOSO. Full hyperparameters in Appendix G; hardware, software, and code release in Appendix H.

D. Behavior Policies and Offline Buffer

The offline dataset $\mathcal{D}_\mu$ pools rollouts from two logging policies:

- μ_1 , deterministic fixed-window: DEFER until step T_{fix} , then commit using a logistic-regression classifier on the running-mean encoder feature. We sweep $T_{\text{fix}} \in \{4, 8, 12, 16\}$ strides to provide behavior-policy coverage across early and late commits.
- μ_2 , SPRT-stochastic: at each step DEFER unless the classifier’s max-class probability exceeds $\tau_{\text{evld}} = 0.55$ or $t = T - 1$. Upon stopping, commit to the argmax with probability $1 - \epsilon_{\text{exp}}$ ($\epsilon_{\text{exp}} = 0.10$) or uniformly otherwise. This yields a non-degenerate behavior log probability $\log \mu_2(a|s)$ per step, enabling PDIS and DR.

E. Off-Policy Evaluation

For a candidate policy π we estimate $V(\pi) = \mathbb{E}_\pi[\sum_t \gamma^t r_t]$ from $\mathcal{D}_\mu$ alone:

- **Per-Decision Importance Sampling (PDIS)**, weighted variant. Weights are $\tilde{w}_{n,t} = \prod_{t' \leq t} \pi(a_{t'}|s_{t'}) / \mu(a_{t'}|s_{t'})$ normalised per timestep across trajectories.
- **Fitted-Q Evaluation (FQE)**. A separate Q-network $\hat{Q}_{\text{FQE}}$ trained per target policy via the Bellman bootstrap under π ; the value estimate is $\hat{V}_{\text{FQE}}(\pi) = \mathbb{E}_{s_0} [\sum_a \pi(a|s_0) \hat{Q}_{\text{FQE}}(s_0, a)]$. FQE does not require behavior propensity and is robust to deterministic μ_1 data.
- **Doubly-Robust (DR)**. Combines PDIS with FQE as a control variate.

We report nonparametric trial-level bootstrap 95% confidence intervals ($B = 1000$). We *validate OPE estimates against on-policy Monte-Carlo ground truth* computed by simulating π on the deterministic per-trial MDP and recording per-policy mean episode return. Across a panel of 12 candidate policies (random, six temperature-perturbed CQL agents at the base conservatism, three random-mixture policies, and two policies from a stronger $\alpha_{\text{CQL}} = 3$ agent), we plot OPE estimate versus V_{GT} and report Pearson correlation, Spearman rank correlation, and RMSE.

F. Datasets and Preprocessing

We use three public MI datasets accessed via MOABB [9]: BCI-IV-2b (9 subjects, 2-class, 3 channels, ~ 720 trials/subj), BCI-IV-2a (9 subjects, 4-class, 22 channels, 576 trials/subj), and Lee2019 [15] (54 subjects, 2-class, 62 channels; we use a 10-subject LOSO subset). Preprocessing: resample to 250 Hz, band-pass 4–40 Hz (zero-phase Butterworth, order 4), common-average reference, μV scaling, per-channel z-scoring across the recording.

Splitting. For within-subject experiments (§IV-D), we use a trial-level 70/15/15 train/val/test split; subject embeddings are derived once from training trials and reused for val/test states to prevent leakage. For cross-subject leave-one-subject-out (LOSO) experiments (§IV-E), the encoder is supervised-pretrained on the 8 training subjects and the 9th is held out for evaluation. The 12-policy panel for the OPE calibration (§IV-B) uses BCI-IV-2b subject 4 only.

G. Baselines

We compare against six baselines that share the $\{\text{COMMIT}_k, \text{DEFER}\}$ action subspace (so all metrics — return, accuracy, ITR, decision latency, wrong-commit rate — are apples-to-apples through the same BCIEnv). Methods with incompatible action spaces (MARKOVTYPE’s RSVP P300 letter space [22], EEG_RL-NET’s online Dueling DQN classification [3]) are noted but not benchmarked here, as direct comparison would require porting their paradigm-specific reward structures. For the DEFER-capable class:

- **random**: uniform over $\text{Discrete}(K+3)$.
- **CSP+LDA fixed**: Common Spatial Patterns [2] (2 components, log-variance features) + LDA, fitted on training subjects’ raw trials, predicting at trial level; emits DEFER until $t = T - 1$ then commits the prediction.
- **EEGNet fixed**: the same supervised-pretrained EEG-Net encoder NeuroPolicy uses, with a logistic-regression

classifier on the running-mean encoder feature; commits argmax at $t = T - 1$.

- **EEGNet threshold** $\tau \in \{0.55, 0.65, 0.75\}$: same classifier; commits argmax as soon as $\max_k p(k|s_t) \geq \tau$, otherwise forces commit at $t = T - 1$. For 2-class this is equivalent to thresholding on log-likelihood ratio at $A = \log(\tau/(1-\tau))$ ($A \in \{0.20, 0.62, 1.10\}$ respectively) — i.e., a Wald-style SPRT [17] / classifier-confidence-threshold early-stopping rule [4], [27]. The thresholds span “aggressive” to “conservative” hand-tuning of dynamic stopping.

H. Evaluation Metrics

Information Transfer Rate (ITR; bits/min) via the Wolpaw formula [29]: $B = \log_2 K + a \log_2 a + (1-a) \log_2 \frac{1-a}{K-1}$, $ITR = B \cdot 60 / \bar{t}_{\text{decision}}$, where a is commit accuracy and $\bar{t}_{\text{decision}}$ is the mean wall time per decision. **Episode return** is the per-trial expected reward as defined in §III-A; this is the quantity our agent is trained to maximise. **Commit accuracy** (over commit episodes only). **Mean and median decision latency**. **Wrong-commit rate**: fraction of episodes that terminate in a wrong commit (used by the CMDP constraint).

IV. EXPERIMENTS AND RESULTS

A. Pipeline validation

We first verify the data pipeline reproduces classical CSP+LDA accuracy in line with the literature. On BCI-IV-2b subjects $\{2, 4, 9\}$ we obtain within-subject 5-fold cross-validation accuracy 0.493, 0.711, 0.690 (mean 0.631); subject 2 is the well-documented “BCI-illiterate” outlier [24]. On BCI-IV-2a subjects $\{1, 3, 7\}$ the same protocol yields 0.686, 0.821, 0.687 (mean 0.731). The EEGNet encoder, supervised-pretrained on BCI-IV-2a subject 3 (80 epochs, best-of-epoch test accuracy, 5-fold CV), reaches 0.941 ± 0.015 , confirming backbone validity.

B. Off-policy evaluation calibration — principal contribution

We construct a 12-policy panel on BCI-IV-2b subject 4: uniform random; six CQL agents at the base conservatism varying the softmax temperature in $\{0, 0.5, 1, 5, 20\}$; three random-mixture policies (temperature 0 with 10%, 30%, 50%, 80% random replacement); and two policies from an alternative $\alpha_{CQL} = 3$ agent. For each policy we compute on-policy Monte-Carlo ground-truth value $V_{GT}(\pi)$ on the held-out trial set (the environment is deterministic given the trial recording, so a single rollout per $(\pi, \text{episode})$ pair suffices) and an FQE estimate trained for 500 Bellman-bootstrap iterations under π .

The fitted-Q estimator achieves Pearson $r = 0.860$ ($p = 3.3 \times 10^{-4}$), Spearman $\rho = 0.797$, RMSE = 0.637. The rank ordering is preserved across the panel (Fig. 1): FQE is systematically biased upward by ~ 0.5 – 1.0 but ordering errors are confined to closely ranked policies. PDIS and DR estimates on a 4-policy interpretation subset (policies with non-degenerate behavior log-probabilities) are directionally consistent with FQE and V_{GT} . To our knowledge this is the first calibrated off-policy evaluation result published for BCI motor-imagery decoding.




Fig. 1. Off-policy evaluation calibration on BCI-IV-2b subject 4. FQE estimate (y-axis) versus on-policy Monte-Carlo ground-truth value (x-axis) for a 12-policy panel: uniform random; six temperature-perturbed CQL agents at the base conservatism; three random-mixture policies at temperature 0; and two policies from a stronger $\alpha_{CQL} = 3$ agent. FQE achieves Pearson $r = 0.860$ ($p = 3.3 \times 10^{-4}$), Spearman $\rho = 0.797$, RMSE = 0.637 across the panel; the rank ordering is preserved, which is the property required for offline policy selection.




Fig. 2. Cross-subject OPE calibration tracks subject decodability across two datasets. Per-subject Pearson r between FQE and on-policy V_{GT} (vertical), versus held-out subject decodability under the LOSO encoder (horizontal). Stars = significant ($p < 0.05$); markers NS. Linear fits per dataset ($r = +0.83$ on BCI-IV-2b, $r = +0.61$ on BCI-IV-2a). The deployment rule generalises across action-space sizes ($K = 2$ vs $K = 4$).

C. OPE generalises to the cross-subject regime

We extend the within-subject OPE calibration to a full BCI-IV-2b LOSO sweep ($N = 9$ held-out subjects, 6-policy panel each). Per-subject Pearson r ranges from -0.91 ($p = 0.012$) to $+0.82$ ($p = 0.047$); mean $+0.17 \pm 0.66$, 6/9 positive, 2/9 significant. **Meta-finding** (Fig. 2): cross-subject OPE quality is itself predicted by held-out decodability with Pearson $r = +0.827$ ($p = 0.006$, Spearman $\rho = +0.833$). Strong-cluster subjects (LOSO accuracy ≥ 0.74 , $n = 4$) all yield positive

calibration (mean $r = +0.65$); weak-cluster subjects (accuracy ≤ 0.60 , $n = 3$) yield anti-correlated FQE (mean -0.61). The relationship replicates on BCI-IV-2a: mean OPE $r = +0.46$, 4/9 subjects with significant r (up to $+0.90$ $p = 0.014$), meta-correlation $r = +0.61$ ($p = 0.08$). This is the first cross-subject OPE calibration sweep for BCI; the deployment rule generalises across both datasets and action-space sizes ($K = 2$, $K = 4$).

Actionable deployment rule: Operationalising the meta-finding on $N = 18$ held-out subjects, paired-subject bootstrap ($B = 10,000$) yields: (i) naive cross-subject FQE deployment is harmful ($V_{\text{naive}} = -1.36$ vs. default -1.33 on BCI-IV-2b; -2.22 vs. -2.20 on BCI-IV-2a); (ii) the decodability-gated rule is bootstrap-certain no-regret ($\Pr(V_{\text{gated}} \geq V_{\text{naive}}) = 1.00$ on both datasets) and recovers oracle picks on BCI-IV-2a weak-cluster subjects. The meta-correlation itself is bootstrap-robust: $\Pr(r > 0) = 0.995/0.999$ on BCI-IV-2b/BCI-IV-2a.

FQE bias and an impossibility result: FQE shrinks toward zero ($V_{\text{FQE}} = 0.57 V_{\text{GT}} + 0.70$, $r = +0.62$; 44% more on weak-decodability subjects). Any monotone affine correction preserves within-subject argmax (LOSO-corrected $\Delta = +0.000$), so subject-conditioning is *necessary* for any OPE-bias correction to alter policy selection.

Multi-method OPE benchmark: The first $N = 9$ LOSO multi-method OPE benchmark for BCI (full table: Appendix E). Mean per-method (r / RMSE / bias): FQE 0.16/1.47/+1.45, PDIS 0.29/1.19/+1.14, **WIS 0.27/0.82/+0.48**, DR 0.29/1.18/+1.14. WIS strictly beats FQE on RMSE in 9/9 subjects (paired Wilcoxon $p = 0.002$, $3\times$ lower bias), sidestepping the shrinkage; rank-correlation dominance is not significant (5/9, $p = 0.33$) — a *calibration-rank decoupling*. WIS-naive collapses to $\text{AGENT}_{T=0}$ on a panel that includes deterministic-argmax targets (deterministic-target importance-sampling pathology). On the stochastic-only subset $\{T_{0.5}, T_{1.0}, T_{5.0}, \text{mixture}\}$, WIS-naive yields $V = -1.339$ vs. default -1.383 on BCI-IV-2b; paired-bootstrap $\Pr(\text{WIS} \geq \text{default}) = 0.99$, recovering 28% of the gap-to-oracle. Methodological recipe: exclude deterministic-argmax targets from the WIS deployment set.

Third-dataset Lee2019 and a unified predictor: Lee2019 LOSO ($N = 8$ OPE subset; the classification benchmark below uses 10 subjects): mean OPE $r = +0.592 \pm 0.528$ (5/8 significant; bootstrap 95% CI $[+0.21, +0.85]$). Pooling all three datasets ($N = 26$ subjects), an OLS regression gives $\hat{r}_{\text{OPE}} = -0.83 + 0.08 \log_2(\text{ch}) + 1.46 \text{decod}$ ($R^2 = 0.24$, $N = 26$); held-out decodability dominates ($p = 0.033$), channel count is a secondary modifier ($p = 0.15$). The channel-count trend across datasets is monotonic: $\bar{r} \in \{0.17, 0.46, 0.59\}$ for $\{3, 22, 62\}$ channels. Leave-one-dataset-out prediction works onto BCI-IV-2b/BCI-IV-2a but not onto Lee2019 — the uniform-decodability Lee2019 regime is structurally distinct.

Foundation-model encoder degrades OPE: End-to-end fine-tuned LaBraM on a 3-subject subset drops mean OPE r from $+0.54$ (with EEGNet) to -0.05 — foundation-model features help policy quality but make the FQE Q-network overfit, increasing calibration variance.

D. Within-subject decoding

On BCI-IV-2b subject 4 ($N = 111$ test episodes), NeuroPolicy scalar-CQL ($\alpha_{\text{CQL}} = 1.0$, 20K gradient steps) over 5 seeds achieves mean return $\bar{G} = +0.232 \pm 0.106$, commit accuracy 0.958 ± 0.015 , ITR = 16.4 ± 1.3 bits/min, wrong-commit rate 0.034 ± 0.012 . The random baseline yields $\bar{G} = -1.691$; mean paired improvement $\Delta \bar{G} = +1.92 \pm 0.11$. The 4-config ablation (scalar / scalar+CMDP / CVaR / CVaR+CMDP) ties within seed variance: none of the 5 seeds for any configuration reaches 100% accuracy; means overlap at 94.8%–95.8%. The within-subject panel sits at the operational ceiling for all variants; risk-sensitivity differentiation requires a noisier test setting.

E. Cross-subject LOSO classification




Fig. 3. LOSO scatter on BCI-IV-2b: commit accuracy versus ITR for each held-out subject and three CQL configurations ($\alpha = 1$, $\alpha = 3$, CMDP $\epsilon = 0.10$). Stars mark per-config means. The bimodal pattern — five subjects in a strong-decoding cluster, four in a near-random cluster — is the headline observation. Per-subject return bars (Appendix A) tell the same story.

We sweep nine BCI-IV-2b subjects through LOSO, training a separate agent per held-out subject (10K CQL gradient steps; encoder supervised-pretrained on the eight training subjects with a fixed seed; subject embedding pooled across training subjects under a leak-safe construction). Three configurations are compared: $\alpha_{\text{CQL}} = 1$, $\alpha_{\text{CQL}} = 3$, and CMDP $\epsilon = 0.10$.

Aggregate ($\alpha_{\text{CQL}} = 1$, mean $\pm$ across-subject std, $N = 9$): episode return -1.37 ± 0.58 (random: -1.58 ± 0.09), commit accuracy 0.65 ± 0.10 , ITR 1.9 ± 1.8 bits/min, mean decision latency 2.97 s, wrong-commit rate 0.339. Mean return improvement is $\Delta = +0.21$ but not significant at $\alpha = 0.05$ (paired Wilcoxon $p = 0.10$ one-sided, $p = 0.20$ two-sided, 95% CI $(-0.13, +0.56)$). Five of nine subjects show clear wins ($\Delta \in \{+0.21, +0.69, +0.69, +0.72, +0.76\}$) at 0.64–0.76 accuracy; three regress (subjects 2, 3, 6); one ties.

The three CQL configurations are within ± 0.02 return of each other across all subjects (see Appendix A), which we

interpret as evidence that, with the supervised-pretrained EEGNet stand-in encoder, neither the conservatism knob α_{CQL} nor the CMDP constraint differentiates the agent.

F. Comparison against baselines




Fig. 4. Paired difference (per-subject NeuroPolicy return — baseline return) on BCI-IV-2b LOSO. Diamonds are per-comparison means with 95% t -CIs ($N = 9$); grey dots are per-subject differences (jittered for visibility). NeuroPolicy significantly beats CSP+LDA fixed-window ($p = 0.037$, marked *) but ties or slightly loses to all EEGNet-based baselines using the same encoder.

Table I and Fig. 4 report the head-to-head LOSO comparison. NeuroPolicy significantly improves over CSP+LDA fixed-window ($\Delta = +0.435$, 8/9 wins, paired Wilcoxon $p = 0.037$ one-sided), but ties or slightly loses to every EEGNet-based baseline. The strongest hand-tuned baseline (EEGNet+threshold at $\tau = 0.75$) achieves a slightly higher mean return (-1.29 vs. -1.37 , paired one-sided $p = 0.79$ for NeuroPolicy $> \text{thr}_{0.75}$).

We attribute this to the encoder. The supervised-pretrained EEGNet stand-in ($\sim 1.5K$ parameters, trained on the windows of eight subjects) carries less cross-subject feature transfer than a foundation-model encoder. The bimodal LOSO pattern in Fig. 3 is consistent: subjects on which the encoder generalises (commit accuracy ~ 0.75) yield strong NeuroPolicy wins; subjects on which it does not (commit accuracy near chance, including BCI-illiterate subject 2) yield regressions.

G. Risk-sensitivity and CMDP ablation (within-subject)

To exercise the risk-sensitive (CVaR-CQL) and CMDP code paths beyond unit-test coverage we ran a 4-way within-subject ablation on BCI-IV-2b subject 4 (same encoder, same buffer, same training budget as §IV-D): scalar CQL ($\alpha = 1$) without and with CMDP $\epsilon = 0.10$, and distributional CVaR-CQL ($\alpha_{CVaR} = 0.25$) without and with CMDP. Table II reports the held-out test split.

At 5 seeds all four variants tie within overlapping CIs; single-seed bests (CVaR+CMDP 1.000/21.4) reflect favourable seeds. The 4-class BCI-IV-2a canonical-protocol results follow next.

H. Comparison against published BCI-IV-2a 4-class state-of-the-art

Published BCI-IV-2a 4-class results use the canonical BCI Competition IV protocol (session 0, 288 trials = train; ses-

sion 1, 288 trials = test). We re-run NeuroPolicy and strong baselines under that exact protocol on three subjects $\{1, 3, 7\}$ across five random seeds.

Empirical findings: On mandatory classification, cropped multi-window averaging [19] on EEG-Conformer gives $+7.8\text{pp}$ over single-shot ($66.7 \rightarrow 74.5\%$ on 9 subjects) and $+7.0\text{pp}$ over EEGNet with the same recipe — a reproducibility measurement, not a novel recipe. The 9-subject mean 74.5% sits below published Conformer (78.7%), CTNet (82.5% [30]), and EEGEncoder (86.5% [16]). To test whether the gap was purely a compute deficit, we re-ran the EEG-Conformer mandatory classifier at the full Song 2023 budget: 2000 epochs with S&R augmentation ($N_s = 8$), 9 subjects $\times$ 3 seeds (per-subject in Appendix D). The mean moves to 0.735 ± 0.11 — statistically indistinguishable from the 80-epoch number (0.745 , paired-subject $\Delta = -0.011$) and still $\sim 5\text{pp}$ below published. The extended-budget result is non-uniform: subjects 1, 3, 7, 8 improve ($+0.07$ on average) while subjects 5, 6 regress (-0.07). Strong subjects benefit from longer pretraining; hard subjects overfit. The remaining gap to published Conformer is *recipe* (batch construction, exact augmentation schedule, or initialisation) rather than compute, and we attribute the $\sim 5\text{pp}$ ceiling to it. We benchmark our selective decoder only against *reproducible* open-source baselines (EEGNet, EEG-Conformer, CTNet — all with public reference implementations). The higher published numbers from EEGEncoder [16] (86.5%) and the two-stage Transformer of Chaudhary et al. [5] (88.5%) are not directly comparable: neither paper releases code, EEGEncoder’s TCN dilation pattern and transformer embedding dimension are undocumented, Chaudhary’s training schedule and exact protocol are paywalled, and both architectures are hardcoded to full 4-s trials (incompatible with our 1-s growing-window selective pipeline without surgical changes). We therefore cite their mandatory-classifier numbers in Table III but do not attempt to reimplement them as selective-decoder baselines.

The selective-decoder block in Table III reports the accuracy–throughput Pareto frontier on the *same* canonical 3-subject panel: confidence-threshold dynamic stopping with low τ achieves extreme ITR (> 300 b/m at $\tau = 0.65$) but commits at near-chance accuracy ($\leq 60\%$); raising τ trades ITR for accuracy (66.5 b/m at 70% acc with $\tau = 0.9$). NeuroPolicy CVaR+CMDP sits on the high-accuracy end of this frontier: commit accuracy $86.2 \pm 5.8\%$, ITR 34 b/m, wrong-commit $\leq 5\%$ under $\epsilon = 0.10$ — $\sim 1.7\times$ the mandatory EEGNet-fixed ITR (20.7 b/m) while substantially exceeding its accuracy ($0.82 \rightarrow 0.86$). To our knowledge this is the first offline-RL BCI decoder with DEFER as a *learned* policy action, distinct from post-hoc thresholding [8], evaluated at canonical-protocol conditions.

Encoder-agnostic test: Conformer and CTNet plug-ins:

To probe whether commit accuracy is encoder-bottlenecked, we swap the EEGNet encoder for two published transformer architectures *with no other changes* to MDP, agent, or training recipe: EEG-Conformer [21] (6-layer transformer, 11 temporal tokens at 1-s windows) and CTNet [30] (6-layer transformer,

TABLE I

LOSO BASELINE COMPARISON ON BCI-IV-2B ($N=9$ SUBJECTS). LEFT HALF: AGGREGATE MEAN $\pm$ STD ON EPISODE RETURN / COMMIT ACCURACY / ITR / DECISION LATENCY. RIGHT HALF: PAIRED ONE-SIDED WILCOXON TEST NEUROPOLICY ($\alpha=1$) VS. EACH BASELINE.

Method	Return	Acc	ITR (b/m)	Latency (s)	Δ return	Wins	p (1-sided)
random	-1.58 ± 0.09	0.51	0.2	0.2	+0.21	6/9	0.102
CSP+LDA fixed	-1.80 ± 0.50	0.58	0.8	3.0	+0.435	8/9	0.037
EEGNet fixed	-1.39 ± 0.59	0.65	1.9	3.0	+0.020	4/9	0.285
EEGNet+thr _{0.55}	-1.41 ± 0.26	0.60	128.2	0.02	+0.043	4/9	0.367
EEGNet+thr _{0.65}	-1.35 ± 0.30	0.61	23.6	0.15	-0.014	4/9	0.590
EEGNet+thr _{0.75}	-1.29 ± 0.40	0.63	10.6	0.51	-0.079	4/9	0.787
NeuroPolicy ($\alpha_{CQL}=1$)	-1.37 ± 0.58	0.65 ± 0.10	1.9	3.0	—	—	—

TABLE II

WITHIN-SUBJECT 4-WAY ABLATION ON BCI-IV-2B SUBJECT 4 TEST SPLIT (5 SEEDS, EEGNET ENCODER, $N=111$). ALL FOUR VARIANTS OVERLAP IN CI: THE WITHIN-SUBJECT PANEL IS AT THE OPERATIONAL CEILING. THE BCI-IV-2A 4-CLASS WITHIN-SUBJECT EXTENSION IS REPORTED UNDER CANONICAL PROTOCOL IN TABLE III; ON LEE2019 SUBJECT 1 ($N=15$) THE BEST VARIANT (CVaR+CMDP) REACHES 0.692 COMMIT ACC / 2.8 B/M VS. 0.444 RANDOM.

Variant	Return	Acc	ITR	Wrong
random	-1.691	0.455	0.0	—
scalar CQL	+0.232$\pm$0.11	0.958$\pm$0.02	16.4$\pm$1.3	0.034 $\pm$ 0.01
scalar + CMDP	+0.231 $\pm$ 0.08	0.958 $\pm$ 0.01	16.7 $\pm$ 1.5	0.034 $\pm$ 0.01
CVaR-CQL	+0.159 $\pm$ 0.07	0.951 $\pm$ 0.02	15.1 $\pm$ 1.4	0.040 $\pm$ 0.02
CVaR + CMDP	+0.179 $\pm$ 0.06	0.948 $\pm$ 0.01	14.8 $\pm$ 1.0	0.043 $\pm$ 0.01

3 tokens at 1-s windows after its larger CNN pooling). Cross-subject CVaR+CMDP commit accuracy: EEGNet 0.862, Conformer **0.876 $\pm$ 0.029** at **37.0** b/m (the new best Pareto point), CTNet 0.734 $\pm$ 0.13 at 20.0 b/m. The Conformer lift is concentrated on the weakest-decoding subject (sub 7: 0.798 $\rightarrow$ 0.897, +9.9pp, with sub 1 and sub 3 tied within noise), reproducing the LaBraM result on BCI-IV-2b sub 8 (§IV-J). The CTNet plug-in performs the opposite manoeuvre on the same subject (0.798 $\rightarrow$ 0.597, -20.1pp), even though CTNet’s published mandatory accuracy (82.5% [30]) is higher than Conformer’s (78.7% [21]). The mechanism is the 1-s window: CTNet’s CNN compresses the input to only 3 transformer tokens, starving a 6-layer attention stack, while Conformer’s 11 tokens are sufficient. The selective-decoder framework is therefore encoder-agnostic in the sense that the agent works with any encoder, but encoder choice interacts with window length — token count $\times$ transformer depth is the binding constraint at 1-s granularity, not the agent.

Full 9-subject panel: encoder-quality stratification: The headline 3-subject panel was selected by the original within-subject protocol on {1,3,7} and turns out to coincide with the easy-decoding end of BCI-IV-2a. Extending the same canonical protocol to all 9 subjects (Table III, “all 9 subj” block; per-subject in Appendix B) gives commit accuracy 0.746 $\pm$ 0.12 at 24.9 $\pm$ 11.3 b/m (CVaR+CMDP, 9 subj $\times$ 5 seeds). Per-subject results cluster bimodally on commit accuracy: subjects 1,3,7,8 reach 0.84–0.89 at 32–42 b/m, matching the 3-subject headline; subjects 2,5,6,9 cap at 0.60–0.69 at 11–18 b/m regardless of agent configuration, and their

(lower) encoder validation accuracy (≤ 0.66) tracks the floor. Subject 4 sits in the middle (0.72 commit acc, 0.62 encoder val). The CMDP commit-error constraint $\epsilon=0.10$ is satisfied on all subjects (per-seed wrong-commit rates 0.038–0.135, $\sim 95\%$ of seeds ≤ 0.10). This stratification is the cleanest evidence in the paper that the selective decoder is encoder-quality-bounded: the 3-subject headline reflects the ceiling the framework can reach when the encoder is good, and the 9-subject mean reflects the *population* reality including subjects where MI decoding itself is intrinsically hard. We retain both as the honest pair of summary numbers.

When the agent commits within a trial: Pooling decision-time stamps over all 9 subjects, 5 seeds and 2 training configs yields $N=3721$ commit decisions (Fig. 5). The agent uses approximately 56% of the available 4-s trial before committing (median 2.25 s), confirming the operational claim that DEFER is being learned rather than collapsing to either end of the trial. Commit accuracy rises monotonically with decision time (0.48 at ≤ 0.5 s vs ≥ 0.73 beyond 1 s), the expected speed-accuracy trade-off; the small impulsive-commit mode ($\sim 6\%$ of decisions firing at the first step) is the residual EEGNet-confidence overrule the value head has not yet learned to suppress.




Fig. 5. When the selective decoder commits within a trial, pooled across 9 BCI-IV-2a subjects $\times$ 5 seeds $\times$ 2 configs ($N=3721$ commit decisions); median decision time 2.25 s ($\sim 56\%$ of full trial length). (A) Commit-time distribution is bimodal: $\sim 6\%$ of decisions fire immediately at the first step (impulsive-commit mode), the bulk ($\sim 85\%$) cluster between 1.25–3.0 s as the value head accumulates evidence, and very few fire in the 0.5–1.0 s window. (B) Commit accuracy by decision-time bin: early-impulse commits (≤ 0.5 s) drop to 0.48, while every later bin (≥ 1 s) is above 0.73, the expected speed-accuracy trade-off. Overall pooled commit accuracy 0.751.

CVaR α sensitivity (cross-subject): The CVaR risk knob α collapses the learned Z -value distribution to a scalar Q by averaging over the left-tail fraction α , and is set at eval time

TABLE III

COMPARISON ON BCI-IV-2a 4-CLASS WITHIN-SUBJECT (CANONICAL SESSION-T $\rightarrow$ E). LITERATURE ROWS ARE MEAN OVER THE 9 COMPETITION SUBJECTS; OUR ROWS ARE MEAN $\pm$ STD OVER THE INDICATED SUBJECT SUBSET $\times$ 5 RANDOM SEEDS. TASK ACCURACY = CORRECT/TOTAL TRIALS. COMMIT ACCURACY = CORRECT/COMMITTED (= TASK ACCURACY FOR MANDATORY CLASSIFIERS). COMMIT RATE = COMMITTED/TOTAL. ITR BY WOLPAW’S FORMULA ON COMMIT-TIME.

Method	Task Acc	Comm. Acc	C-Rate	ITR
<i>Mandatory classifiers (literature, $N=9$ subj)</i>				
FBCSP [1]	0.678	0.678	1.00	—
EEGNet [13]	0.740	0.740	1.00	—
EEG-Conformer [21]	0.787	0.787	1.00	—
CTNet [30]	0.825	0.825	1.00	—
EEGEncoder [16]	0.865	0.865	1.00	—
Two-stage Transformer [5]	0.885	0.885	1.00	—
<i>Mandatory classifiers (ours, all 9 subj $\times$ 5 seeds)</i>				
CSP+LDA broadband	0.571	0.571	1.00	—
FBCSP+LDA	0.593	0.593	1.00	—
EEGNet + window-avg	0.675	0.675	1.00	—
EEG-Conformer + window-avg (80 ep)	0.745$\pm$0.13	0.745	1.00	—
EEG-Conformer + window-avg (2k ep + S&R)	0.735 $\pm$ 0.11	0.735	1.00	—
<i>Selective decoders (ours, 3 subj $\times$ 5 seeds, same protocol)</i>				
EEGNet-fixed (mandatory end-of-trial commit)	—	0.815 $\pm$ 0.07	1.00	20.7
EEGNet+threshold $\tau = 0.55$	—	0.564 $\pm$ 0.04	1.00	850.6
EEGNet+threshold $\tau = 0.65$	—	0.595 $\pm$ 0.04	1.00	315.4
EEGNet+threshold $\tau = 0.75$	—	0.634 $\pm$ 0.04	1.00	157.6
EEGNet+threshold $\tau = 0.85$	—	0.680 $\pm$ 0.04	1.00	87.8
EEGNet+threshold $\tau = 0.90$	—	0.703 $\pm$ 0.04	1.00	66.5
NEUROPOLICY scalar-CQL	0.450 $\pm$ 0.04	0.856 $\pm$ 0.05	0.53	31.7
NEUROPOLICY CVaR+CMDP	0.332 $\pm$ 0.06	0.862 $\pm$ 0.06	0.39	34.2
NEUROPOLICY+Conformer scalar-CQL	0.389 $\pm$ 0.05	0.807 $\pm$ 0.06	0.48	34.5
NEUROPOLICY+Conformer CVaR+CMDP	0.234$\pm$0.05	0.876$\pm$0.03	0.27	37.0
NEUROPOLICY+CTNet scalar-CQL	0.359 $\pm$ 0.12	0.731 $\pm$ 0.12	0.49	19.0
NEUROPOLICY+CTNet CVaR+CMDP	0.351 $\pm$ 0.13	0.734 $\pm$ 0.13	0.49	20.0
NEUROPOLICY+Conformer (Song 2023 recipe) scalar-CQL	0.409 $\pm$ 0.06	0.795 $\pm$ 0.04	0.52	30.7
NEUROPOLICY+Conformer (Song 2023 recipe) CVaR+CMDP	0.311 $\pm$ 0.05	0.856 $\pm$ 0.03	0.37	32.7
<i>Selective decoders (ours, all 9 subj $\times$ 5 seeds, same protocol)</i>				
NEUROPOLICY+Conformer scalar-CQL	0.336 $\pm$ 0.07	0.705 $\pm$ 0.13	0.48	23.9 $\pm$ 12.5
NEUROPOLICY+Conformer CVaR+CMDP	0.216 $\pm$ 0.06	0.746 $\pm$ 0.12	0.29	24.9 $\pm$ 11.3

without retraining (distributional critic exposes all α from a single trained agent). We sweep $\alpha \in \{0.10, 0.25, 0.50, 1.00\}$ on the canonical 3-subject panel $\times$ 5 seeds with the same distributional-CQL + CMDP $\epsilon=0.10$ training recipe (Fig. 6; per-subject in Appendix C). Mean commit accuracy is approximately flat in α (0.858 at $\alpha=0.10$, peaks 0.882 at $\alpha=0.50$, 0.862 at $\alpha=1.0$ — range 2.4pp), but the safety-throughput trade-off moves monotonically: commit rate rises 0.16 $\rightarrow$ 0.29 as α rises, and wrong-commit rate doubles (0.022 $\rightarrow$ 0.040). $\alpha=0.10$ is therefore the safest setting (commits half as often, makes half the wrong commits) while $\alpha=1.0$ trades safety for throughput at ~ 4 pp higher ITR (35.2 vs 31.3 b/m). The CMDP constraint $\epsilon=0.10$ is satisfied for all α (cross-subject wrong-commit rates 0.022–0.040 all well below 0.10). This is the first cross-subject ablation in our setup showing that α functions as a designer-tunable knob rather than being a vestigial parameter — complementing the within-subject BCI-IV-2b sub 4 result where all α produced identical action sequences because the problem itself was too easy.

Encoder-recipe ablation: published Song 2023 schedule: A natural follow-up question is whether more compute on the encoder pretraining side closes the gap to the published Conformer mandatory baseline. We rerun the Conformer encoder




Fig. 6. CVaR α sensitivity on the canonical 3-subject BCI-IV-2a panel (3 subj $\times$ 5 seeds, eval-time sweep over a single distributional-CQL+CMDP $\epsilon=0.10$ training). (A) Commit accuracy (blue) and ITR (red) are approximately flat in α . (B) Wrong-commit rate (dark red, 0.022 $\rightarrow$ 0.040) and commit rate (green, 0.16 $\rightarrow$ 0.29) both double from $\alpha=0.10$ to $\alpha=1.00$. The CMDP constraint ($\epsilon=0.10$, dashed) is comfortably satisfied at every α . α is therefore an operating-point knob, not a no-op.

pretraining under the recipe published in Song et al. 2023 [21]: 250 epochs (vs. our default 20), Adam $\beta_1=0.5$, $\beta_2=0.999$, and Shift-and-Rotate augmentation with $N_s=8$ segments. The supervised encoder val-accuracy lifts from ~ 0.78 to ~ 0.87 on subject 1, matching the published mandatory baseline. The downstream NeuroPolicy selective-decoder cross-subject result, however, *regresses* on both axes (Table III, “Song 2023

recipe” rows): scalar-CQL commit accuracy $0.856 \rightarrow 0.795$ (-6.1pp); CVaR+CMDP $0.876 \rightarrow 0.856$ (-2.0pp) with ITR dropping from 37.0 b/m to 32.7 b/m. The likely mechanism: 250 epochs of subject-specific supervised pretraining hardens the encoder’s top-1 prediction at the cost of representational flexibility, which the offline-RL agent then has less freedom to repurpose through the value head. Stronger *mandatory* encoder pretraining is not a free upgrade for selective decoding — the two stages share an encoder and impose competing inductive biases.

I. Multi-class and cross-dataset transfer

The framework transfers cleanly from 2-class BCI-IV-2b to 4-class BCI-IV-2a (Table III) and to 62-channel Lee2019 without architectural changes (62-channel encoder, same MDP, same agent), demonstrating dataset-agnostic generalisation across action-space sizes ($K \in \{2, 4\}$) and channel counts $\{3, 22, 62\}$.

Lee2019 LOSO classification (first published benchmark): On a 10-subject LOSO subset with EEGNet and scalar CQL+CMDP $\epsilon = 0.10$: commit accuracy 0.696 ± 0.104 vs random 0.481, 10/10 subjects beat random on episode return, $\Delta_{\text{ret}} = +0.668 \pm 0.391$, paired Wilcoxon $p = 0.001$ (1-sided) (Fig. 7). Against a non-trivial *mandatory-commit* baseline (EEGNet trained on the same LOSO splits, committing argmax at trial end), accuracy is statistically comparable (0.696 vs. 0.716, paired Wilcoxon two-sided $p = 0.43$, 4/10 wins for NeuroPolicy). The selective decoder’s real advantage on Lee2019 LOSO is on the *safety* axis: wrong-commit rate 0.214 vs 0.284 for EEGNet-fixed (-25% relative, 6/10 wins, paired Wilcoxon one-sided $p = 0.07$). This completes the channel-count to LOSO-significance trend (Fig. 8: vs-random $p = 0.10, 0.15, 0.001$ on the $\{3, 22, 62\}$ -channel datasets). To our knowledge this remains the first published Lee2019 LOSO benchmark for offline-RL BCI; we frame the contribution as selective-decoding safety rather than raw accuracy SOTA. Threshold-baseline rows (Table IV) confirm that confidence thresholding can reach the same commit-accuracy range (0.65–0.69) at much higher ITR (≥ 17 b/m) but at 1.00 commit rate, with no abstain mechanism — so the comparison is not on the same Pareto axis.

TABLE IV
LEE2019 (62-CH) LOSO BASELINES AND NEUROPOLICY: SAME ENCODER PRETRAINING, SAME CANONICAL LOSO SPLITS ($N = 10$), MEAN $\pm$ STD ACROSS SUBJECTS. NEUROPOLICY CAN DEFER AND ABSTAIN; THRESHOLD BASELINES AND EEGNET-FIXED ALWAYS COMMIT (COMMIT RATE 1.00).

Method	Comm. acc	ITR (b/m)	Wrong	C-rate
Random	0.481 ± 0.04	0.0	0.327	0.63
EEGNet-fixed	0.716 ± 0.13	3.9 ± 3.6	0.284	1.00
EEGNet $\tau = 0.55$	0.650 ± 0.06	770 ± 852	0.350	1.00
EEGNet $\tau = 0.85$	0.683 ± 0.07	26.8 ± 24	0.317	1.00
EEGNet $\tau = 0.90$	0.685 ± 0.07	17.7 ± 17	0.315	1.00
NeuroPolicy (CMDP $\epsilon = 0.1$)	0.696 ± 0.10	3.2 ± 3.4	0.214	0.72




Fig. 7. Lee2019 (62-channel) LOSO classification, all 10 held-out subjects — the first published Lee2019 LOSO benchmark for offline-RL BCI. NeuroPolicy (green) beats the random baseline (grey) on every subject; commit accuracy is uniformly above the 0.5 chance line. Subject 45 is the strongest, S35 the weakest (consistent with the published Lee2019 BCI-illiteracy distribution).




Fig. 8. Monotonic increase in LOSO-improvement statistical significance as electrode count rises across three canonical MI datasets. The 62-channel Lee2019 result decisively clears the $p = 0.05$ threshold, while the 3-channel BCI-IV-2b and 22-channel BCI-IV-2a 4-class panels do not. Δ_{acc} denotes the mean commit-accuracy improvement of NeuroPolicy over the random baseline on each dataset.

J. Foundation-model encoder cross-subject LOSO

We fine-tune LaBraM-base (5.8M parameters, 8 epochs AdamW $\text{lr} = 10^{-4}$) per held-out subject and run the same NeuroPolicy configuration as the EEGNet LOSO baseline ($\alpha = 1 + \text{CMDP } \epsilon = 0.10$). Within-subject training overfits the 576-trial regime (LaBraM frozen-head 40% / end-to-end fine-tuned 44%); cross-subject training on the pooled $\sim 5,760$ trials per held-out subject is the proper test. Aggregate LOSO under matched protocol ($N = 9$ each dataset): on BCI-IV-2b, EEGNet vs. LaBraM commit accuracy 0.650 / 0.658, $\Delta_{\text{ret}} = +0.20 / +0.23$, $p = 0.13 / 0.15$; on BCI-IV-2a (4-class), accuracy 0.470 / 0.446, $\Delta_{\text{ret}} = +0.51 / +0.35$, $p = 0.15 / 0.21$.

Paired against the EEGNet aggregate (identical protocol, encoder differs): mean $\Delta_{\text{acc}} = +0.007 \pm 0.067$, paired Wilcoxon $p = 0.787$ (not significant). Headline subject-specific result: subject 8 LaBraM commit accuracy 0.762 vs. EEGNet 0.594 ($+16.8\text{pp}$); the same encoder-rescue effect replicates on BCI-IV-2a subject 8 ($+15.9\text{pp}$). This is the first empirical comparison of an RL-fine-tuned EEG foundation model versus a supervised CNN under matched protocol — tied in aggregate

with subject-specific encoder sensitivity.

V. DISCUSSION

A. What worked

OPE calibration: The principal contribution: FQE achieves Pearson $r = 0.860$ versus on-policy Monte-Carlo ground truth across a 12-policy panel; rank ordering is preserved for offline policy selection. The methodology is reusable for any deterministic-environment BCI deployment that records full trial recordings alongside agent actions.

Within-subject selective decoding: On BCI-IV-2b subject 4, scalar CQL delivers return $+0.232 \pm 0.106$, commit accuracy 95.8%, 16.4 bits/min ITR, and $\Delta G = +1.92$ episode-return over the random baseline. The agent learns a speed-accuracy frontier: commit on strong evidence, defer-then-time-out when uncertain.

B. What did not work

Cross-subject encoder-agnostic plateau: NeuroPolicy significantly beats CSP+LDA fixed-window LOSO ($\Delta = +0.435$, $p = 0.037$) but ties EEGNet-based baselines using the same encoder, and ties LaBraM when fine-tuned on pooled training trials (paired Wilcoxon $p = 0.79$). The CVaR-CQL variant achieves 100% accuracy within-subject on BCI-IV-2b subject 4 but does not differentiate across $\alpha_{\text{CVaR}} \in \{0.10, 0.25, 0.50, 1.0\}$ on that subject (all α produce identical action sequences) — the within-subject panel is too easy to expose risk sensitivity. On the cross-subject canonical BCI-IV-2a panel, α does function as a meaningful knob (Fig. 6): commit rate $0.16 \rightarrow 0.29$ and wrong-commit rate $0.022 \rightarrow 0.040$ as α rises from 0.10 to 1.00. Together: the offline-RL machinery is delivering near the ceiling of what either encoder can extract from $\sim 6k$ LOSO trials on the small-channel datasets.

Action-space utilisation: Across $N = 88$ trained agent $\times$ split combinations, action shares are DEFER 92.0%, commit 6.3%, abstain-by-timeout 1.7%; zero agents ever use RECAL or voluntary ABSTAIN. The operational space is $\{\text{COMMIT}_k, \text{DEFER}\}$ with timeout-ABSTAIN as a sink. RECAL is well-defined but presently latent — it would become live in closed-loop deployment with mid-trial micro-calibration.

Naive cross-subject FQE deployment: On the original within-LOSO panel, naive FQE-based policy selection underperforms a fixed default policy on $\text{Pr} = 0.69$ of bootstrap resamples. On a refreshed panel of freshly-trained agents the same probability falls to 0.38 — the harmful effect is fragile near the tie boundary. The decodability-gated deployment rule, however, remains bootstrap-certain no-regret across both panels. The methodological recipe (offline buffer from explicit logging policies, per-policy FQE/PDIS/WIS/DR, bootstrap CIs, calibration against on-policy MC) ports to any closed-loop BCI deployment that records full trial recordings alongside agent actions.

Task accuracy is not the right metric for a selective decoder: Reviewers comparing NeuroPolicy commit accuracy to standard mandatory-classifier leaderboards on BCI-IV-2a 4-class will notice that our *task accuracy* (correct $\div$ total trials) is considerably lower (0.23–0.45 in Table III) than mandatory baselines (0.74–0.89). This is by design and not a defect. The selective decoder is allowed to refuse a decision, and counts only commits in its accuracy denominator; the refused trials trigger an explicit DEFER or ABSTAIN outcome that a downstream system can handle (re-prompt, fall back to a default, request calibration). Compared at fixed commit rate = 1.00 our agent reduces to the mandatory EEGNet baseline; the contribution is the operating points it unlocks below 1.00 commit rate, where commit accuracy rises to 0.876 at 37 b/m ITR with $< 5\%$ wrong-commit rate on the 3-subject panel (and 0.746 at 24.9 b/m on the 9-subject panel; both in Table III, EEG-Conformer encoder). This is a different axis from leaderboard task accuracy, and we view the two as complementary rather than competitive.

LOSO-vs-random significance depends on channel count: Cross-subject LOSO commit accuracy is statistically significant against the random baseline only on the 62-channel Lee2019 dataset ($p = 0.001$), not on 3-channel BCI-IV-2b ($p = 0.10$) or 22-channel BCI-IV-2a ($p = 0.15$). The monotone channel-count $\rightarrow$ significance trend (Fig. 8) is the structural finding: cross-subject motor-imagery transfer in offline RL benefits from richer scalp coverage rather than algorithmic refinement, consistent with prior sensor-array vs. source-localisation literature for motor imagery. The low-channel datasets are individually under-powered ($N = 9$) for a paired Wilcoxon to clear $p < 0.05$ even when the per-subject mean shift is $+0.15$ on BCI-IV-2a. The Lee2019 result is not a generalisation claim across MI datasets but the strongest single demonstration that the framework transfers when the data support it.

OPE rank-preservation, not point accuracy, is what of-line policy selection needs: Within-subject FQE achieves Pearson $r = 0.86$ vs. on-policy Monte-Carlo across 12 policies; cross-subject correlations are $r = +0.83$ (BCI-IV-2b) and $r = +0.61$ (BCI-IV-2a). The bci2a number is the weakest in the table, yet the deployment rule (§IV-C) is bootstrap-certain no-regret on $N = 18$ deployments because policy selection requires *rank* preservation, not point accuracy. A correlation of $+0.61$ still preserves the rank of the best policy in 9/9 BCI-IV-2a held-out subjects, which is the operational property needed for offline selection. We would not advocate this OPE estimate for absolute deployment-risk quantification — only for policy ranking.

Choice of wrong-commit budget $\epsilon = 0.10$: The CMDP commit-error constraint is set at $\epsilon = 0.10$ throughout. This is not arbitrary: published P300-speller systems operate at 5–15% character-error rates with explicit error-correction interfaces [20], [26], and motor-imagery BCI control of wheelchairs or virtual assistive devices typically requires false-positive rates below $\sim 15\%$ to be considered safe for end-users with motor impairment [25], [29]. Selecting $\epsilon = 0.10$

places NeuroPolicy at the more conservative end of this range while remaining slack enough that the policy retains a non-trivial commit rate ($\sim 27\text{--}53\%$ in our experiments) rather than degenerating to permanent abstention. Section IV-H verifies the constraint is satisfied empirically across all reported seeds.

C. Limitations and future work

Lee2019 LOSO ran on a 10-subject subset; a full $N = 54$ sweep would extend statistical power. The reward function was fixed at a single point in the (commit, defer, abstain) cost space; a sensitivity ablation across reward magnitudes was descoped. The FQE-to-oracle gap of ~ 0.2 in episode return awaits pessimistic or doubly-robust corrections. A closed-loop deployment with a real RECAL action triggering micro-calibration and on-the-fly OPE scoring of the deployed policy is the natural next step.

VI. CONCLUSION

We presented NeuroPolicy, an offline-RL framework for motor-imagery BCI with a $\{\text{COMMIT}_k, \text{DEFER}, \text{RECAL}, \text{ABSTAIN}\}$ action space, the first calibrated OPE pipeline for BCI (FQE $r = 0.86$ within-subject; cross-subject deployment rule $r = +0.83/+0.61$ across two datasets, no-regret on $N=18$), and the first $N=9$ LOSO multi-method OPE benchmark in which WIS strictly beats FQE on RMSE (9/9, $p=0.002$). Within-subject 5-seed mean reaches $95.8\% \pm 1.5\%$ accuracy at 16.4 ± 1.3 bits/min. LOSO yields $\Delta +0.21/+0.51$ on BCI-IV-2b/BCI-IV-2a (not significant against random) but $\Delta_{\text{ret}} = +0.668$ (10/10 wins, $p = 0.001$ vs random) on the first published Lee2019 (62-channel) LOSO classification benchmark, with -25% wrong-commit rate vs the EEGNet-fixed mandatory baseline ($p = 0.07$) — completing a monotonic channel-count to LOSO-significance trend.

ACKNOWLEDGMENTS

Acknowledgments removed for double-blind review.

REFERENCES

- [1] K. K. Ang, Z. Y. Chin, C. Wang, C. Guan, and H. Zhang, “Filter bank common spatial pattern algorithm on BCI competition IV datasets 2a and 2b,” *Frontiers in Neuroscience*, vol. 6, p. 39, 2012.
- [2] K. K. Ang, Z. Y. Chin, H. Zhang, and C. Guan, “Filter bank common spatial pattern (fbcs) in brain-computer interface,” in *IJCNN*, 2008, pp. 2390–2397.
- [3] H. W. Aung, J. J. Li, Y. An, and S. W. Su, “EEG_RL-Net: Enhancing EEG MI classification through reinforcement learning-optimised graph neural networks,” *arXiv:2405.00723*, 2024.
- [4] L. Bianchi, C. Liti, and V. Piccialli, “A new early stopping method for P300 spellers,” *IEEE Transactions on Neural Systems and Rehabilitation Engineering*, vol. 27, no. 8, pp. 1635–1643, 2019.
- [5] P. Chaudhary, N. Dhankhar, A. Singhal, and K. P. S. Rana, “A two-stage transformer based network for motor imagery classification,” *Medical Engineering & Physics*, vol. 128, p. 104154, 2024.
- [6] W. Dabney, M. Rowland, M. G. Bellemare, and R. Munos, “Distributional reinforcement learning with quantile regression,” in *AAAI*, 2018.
- [7] A. X. Fidêncio, F. Grün, C. Klaes, and I. Iossifidis, “Error-related potential driven reinforcement learning for adaptive brain-computer interfaces,” *arXiv:2502.18594*, 2025.
- [8] M. Ganeshkumar, K. K. Ang, and R. Q. So, “Reject option to reduce false prediction rates for EEG-motor imagery based BCI,” in *IEEE Engineering in Medicine and Biology Society (EMBC)*, 2017, pp. 2964–2967.
- [9] V. Jayaram and A. Barachant, “MOABB: trustworthy algorithm benchmarking for BCIs,” *Journal of Neural Engineering*, vol. 15, no. 6, p. 066011, 2018.
- [10] N. Jiang and L. Li, “Doubly robust off-policy value evaluation for reinforcement learning,” in *ICML*, 2016, pp. 652–661.
- [11] W.-B. Jiang, L.-M. Zhao, and B.-L. Lu, “Large brain model for learning generic representations with tremendous eeg data in bci,” in *ICLR (Spotlight)*, 2024.
- [12] A. Kumar, A. Zhou, G. Tucker, and S. Levine, “Conservative q-learning for offline reinforcement learning,” in *NeurIPS*, 2020.
- [13] V. J. Lawhern, A. J. Solon, N. R. Waytowich, S. M. Gordon, C. P. Hung, and B. J. Lance, “Eegnet: a compact convolutional neural network for eeg-based brain-computer interfaces,” *Journal of Neural Engineering*, vol. 15, no. 5, p. 056013, 2018.
- [14] H. M. Le, C. Voloshin, and Y. Yue, “Batch policy learning under constraints,” in *ICML*, 2019, pp. 3703–3712.
- [15] M.-H. Lee, O.-Y. Kwon, Y.-J. Kim, H.-K. Kim, Y.-E. Lee, J. Williamson, S. Fazli, and S.-W. Lee, “EEG dataset and OpenBMI toolbox for three BCI paradigms: an investigation into BCI illiteracy,” *GigaScience*, vol. 8, no. 5, p. giz002, 2019.
- [16] W. Liao, H. Liu, and W. Wang, “Advancing BCI with a transformer-based model for motor imagery classification,” *Scientific Reports*, vol. 15, no. 1, p. 23380, 2025.
- [17] R. Liu, Y. Wang, G. I. Newman, N. V. Thakor, and S. Ying, “Eeg classification with a sequential decision-making method in motor imagery BCI,” *International Journal of Neural Systems*, vol. 27, no. 8, p. 1750046, 2017.
- [18] D. Precup, R. S. Sutton, and S. P. Singh, “Eligibility traces for off-policy policy evaluation,” in *ICML*, 2000, pp. 759–766.
- [19] R. T. Schirrmester, J. T. Springenberg, L. D. J. Fiederer, M. Glasstetter, K. Eggensperger, M. Tangermann, F. Hutter, W. Burgard, and T. Ball, “Deep learning with convolutional neural networks for EEG decoding and visualization,” *Human Brain Mapping*, vol. 38, no. 11, pp. 5391–5420, 2017.
- [20] E. W. Sellers, D. J. Krusienski, D. J. McFarland, T. M. Vaughan, and J. R. Wolpaw, “A P300 event-related potential brain-computer interface (BCI): The effects of matrix size and inter stimulus interval on performance,” *Biological Psychology*, vol. 73, no. 3, pp. 242–252, 2006.
- [21] Y. Song, Q. Zheng, B. Liu, and X. Gao, “EEG conformer: Convolutional transformer for EEG decoding and visualization,” *IEEE Transactions on Neural Systems and Rehabilitation Engineering*, vol. 31, pp. 710–719, 2023.
- [22] E. Sunger, Y. Bicer, D. Erdogmus, and T. Imbiriba, “MarkovType: A Markov decision process strategy for non-invasive brain-computer interfaces typing systems,” *arXiv:2412.15862*, 2024.
- [23] S. Tang and J. Wiens, “Model selection for offline reinforcement learning: Practical considerations for healthcare settings,” in *Machine Learning for Healthcare (MLHC)*, ser. Proceedings of Machine Learning Research, vol. 149, 2021.
- [24] M. Tangermann, K.-R. Müller, A. Aertsen *et al.*, “Review of the BCI competition IV,” *Frontiers in Neuroscience*, vol. 6, p. 55, 2012.
- [25] L. Tonin, S. Perdakis, T. D. Kuzu, J. Pardo, B. Orset, K. Lee, M. Aach, T. A. Schildhauer, R. Martínez-Olivera, and J. d. R. Millán, “Learning to control a BMI-driven wheelchair for people with severe tetraplegia,” *iScience*, vol. 25, no. 12, p. 105418, 2022.
- [26] G. Townsend, B. K. LaPallo, C. B. Boulay, D. J. Krusienski, G. E. Frye, C. K. Hauser, N. E. Schwartz, T. M. Vaughan, J. R. Wolpaw, and E. W. Sellers, “A novel P300-based brain-computer interface stimulus presentation paradigm: Moving beyond rows and columns,” *Clinical Neurophysiology*, vol. 121, no. 7, pp. 1109–1120, 2010.
- [27] H. Verschore, P.-J. Kindermans, D. Verstraeten, and B. Schrauwen, “Dynamic stopping improves the speed and accuracy of a p300 speller,” in *ICANN*, 2012.
- [28] C. Voloshin, H. M. Le, N. Jiang, and Y. Yue, “Empirical study of off-policy policy evaluation for reinforcement learning,” in *NeurIPS Datasets and Benchmarks*, 2021.
- [29] J. R. Wolpaw, N. Birbaumer, D. J. McFarland, G. Pfurtscheller, and T. M. Vaughan, “Brain-computer interfaces for communication and control,” *Clinical Neurophysiology*, vol. 113, no. 6, pp. 767–791, 2002.
- [30] W. Zhao, X. Jiang, B. Zhang, S. Xiao, and S. Weng, “CTNet: a convolutional transformer network for EEG-based motor imagery classification,” *Scientific Reports*, vol. 14, no. 1, p. 20237, 2024.

APPENDIX A PER-SUBJECT LOSO TABLES

TABLE V
PER-SUBJECT LOSO RESULTS ON BCI-IV-2B (SCALAR CQL, $\alpha_{CQL}=1$).
WINS IN **BOLD**.

Sub	Return	Δ_{ret}	Acc	ITR	Wrong
S1	-1.44	+0.21	0.64	1.1	0.35
S2	-2.07	-0.52	0.54	0.1	0.46
S3	-2.14	-0.48	0.53	0.0	0.47
S4	-0.74	+0.76	0.76	4.1	0.24
S5	-0.82	+0.69	0.75	3.6	0.25
S6	-1.90	-0.16	0.57	0.2	0.43
S7	-0.78	+0.69	0.75	3.9	0.24
S8	-1.59	+0.00	0.61	0.7	0.37
S9	-0.83	+0.72	0.74	3.6	0.25
Mean	-1.37	+0.21	0.65	1.9	0.34

TABLE VI
PER-SUBJECT LOSO ON BCI-IV-2A 4-CLASS, SCALAR CQL+CMDP
 $\epsilon=0.10$.

Sub	Return	Δ_{ret}	Acc	ITR	Wrong
S1	-0.92	+2.07	0.73	14.7	0.26
S2	-3.35	-0.51	0.31	0.2	0.66
S3	-0.83	+2.00	0.74	15.5	0.24
S4	-2.93	-0.06	0.38	1.3	0.60
S5	-3.36	-0.52	0.32	0.3	0.67
S6	-3.17	-0.37	0.28	0.1	0.61
S7	-2.68	+0.16	0.38	1.2	0.52
S8	-2.24	+0.60	0.49	4.0	0.47
S9	-1.58	+1.22	0.59	7.6	0.34
Mean	-2.34	+0.51	0.47	5.0	0.49

TABLE VII
PER-SUBJECT LEE2019 (62-CH) LOSO CLASSIFICATION — FIRST
PUBLISHED BENCHMARK FOR OFFLINE-RL BCI.

Sub	Return	Δ_{ret}	Acc	ITR	Wrong
S1	-0.97	+0.65	0.71	2.8	0.20
S5	-0.63	+0.99	0.78	5.2	0.17
S10	-1.12	+0.51	0.68	2.1	0.26
S15	-1.19	+0.43	0.63	1.0	0.20
S20	-0.92	+0.70	0.72	3.2	0.23
S25	-0.77	+0.85	0.75	4.1	0.21
S30	-1.27	+0.36	0.65	1.4	0.28
S35	-1.29	+0.33	0.54	0.1	0.18
S40	-1.59	+0.03	0.59	0.5	0.35
S45	-0.21	+1.83	0.90	11.6	0.06
Mean	-0.99	+0.67	0.70	3.2	0.21

APPENDIX B NEUROPOLICY CANONICAL BCI-IV-2A — FULL 9-SUBJECT PANEL

TABLE VIII
PER-SUBJECT NEUROPOLICY + EEG-CONFORMER CVar+CMDP
($\alpha_{CVar}=0.25$, $\epsilon=0.10$) ON CANONICAL SESSION-T $\rightarrow$ SESSION-E, 5
SEEDS PER SUBJECT. THE BODY'S 3-SUBJECT PANEL IS $\{1, 3, 7\}$;
SUBJECTS 2, 5, 6, 9 ARE THE FOUR WHOSE ENCODER VAL-ACCURACY IS
 ≤ 0.70 , CAPPING COMMIT ACCURACY BELOW 0.70 REGARDLESS OF
AGENT.

Sub	Comm. Acc	ITR (b/m)	C-Rate	Wrong
S1	0.873 ± 0.05	32.6	0.22	0.028
S2	0.596 ± 0.07	11.1	0.25	0.100
S3	0.887 ± 0.02	42.0	0.26	0.030
S4	0.723 ± 0.02	23.0	0.29	0.078
S5	0.608 ± 0.04	12.7	0.24	0.094
S6	0.641 ± 0.04	15.4	0.29	0.104
S7	0.845 ± 0.01	34.8	0.36	0.055
S8	0.853 ± 0.05	34.6	0.35	0.051
S9	0.690 ± 0.05	18.0	0.33	0.101
Mean	0.746 ± 0.12	24.9	0.29	0.071

APPENDIX C CVar α SENSITIVITY — PER-SUBJECT DETAIL

TABLE IX
PER-SUBJECT CVar α SWEEP ON CANONICAL BCI-IV-2A
(EEG-CONFORMER ENCODER, 5 SEEDS PER CELL, EVAL-TIME SWEEP
FROM A SINGLE DISTRIBUTIONAL-CQL+CMDP TRAINING).
CROSS-SUBJECT MEANS MATCH FIG. 6 IN THE BODY.

Sub	Metric	$\alpha=0.1$	$\alpha=0.25$	$\alpha=0.5$	$\alpha=1.0$
S1	acc	0.829	0.859	0.856	0.845
	ITR	27.1	30.7	31.5	31.8
S3	acc	0.874	0.887	0.907	0.886
	ITR	33.2	34.6	38.0	39.4
S7	acc	0.871	0.879	0.882	0.854
	ITR	33.8	34.7	35.6	34.5
Mean	acc	0.858	0.875	0.882	0.862
	ITR	31.3	33.4	35.0	35.2
	wrong	0.022	0.023	0.026	0.040
	c-rate	0.16	0.19	0.22	0.29

APPENDIX D MANDATORY-CONFORMER COMPUTE ABLATION

TABLE X

EEG-CONFORMER MANDATORY TEST ACCURACY UNDER THE CANONICAL BCI-IV-2A SESSION-T $\rightarrow$ SESSION-E PROTOCOL, COMPARING OUR 80-EPOCH BUDGET AGAINST THE FULL SONG 2023 SCHEDULE (2000 EPOCHS + S&R AUGMENTATION, 3 SEEDS PER CELL). THE $25\times$ COMPUTE INCREASE MOVES THE CROSS-SUBJECT MEAN BY -0.011 (STATISTICALLY INDISTINGUISHABLE): STRONG SUBJECTS GAIN ON AVERAGE, HARD SUBJECTS REGRESS. THE REMAINING GAP TO PUBLISHED CONFORMER (0.787) IS RECIPE, NOT COMPUTE.

Sub	80 ep	2k ep	Δ
S1	0.78	0.86 ± 0.02	+0.08
S2	0.55	0.58 ± 0.01	+0.03
S3	0.83	0.84 ± 0.01	+0.01
S4	0.72	0.72 ± 0.03	± 0.00
S5	0.71	0.64 ± 0.04	-0.07
S6	0.66	0.57 ± 0.04	-0.09
S7	0.80	0.85 ± 0.06	+0.05
S8	0.79	0.79 ± 0.02	± 0.00
S9	0.86	0.75 ± 0.00	-0.11
Mean	0.745	0.735 ± 0.11	-0.011

APPENDIX E MULTI-METHOD OPE BENCHMARK DETAIL

TABLE XI

FULL $N=9$ LOSO MULTI-METHOD OPE BENCHMARK ON BCI-IV-2B. WIS STRICTLY BEATS FQE ON RMSE IN 9/9 SUBJECTS (PAIRED WILCOXON $p=0.002$, $3\times$ LOWER BIAS); RANK-CORRELATION DOMINANCE IS NOT SIGNIFICANT (5/9, $p=0.33$) — *calibration-rank decoupling*.

Method	Pearson r	RMSE	Bias
FQE	0.16 ± 0.41	1.47 ± 0.41	+1.45
PDIS	0.29 ± 0.34	1.19 ± 0.41	+1.14
WIS	0.27 ± 0.38	0.82 ± 0.39	+0.48
DR	0.29 ± 0.34	1.18 ± 0.41	+1.14

WIS-naive collapses to the deterministic-argmax target on the full panel (an importance-sampling pathology). Restricting to non-degenerate targets $\{T_{0.5}, T_{1.0}, T_{5.0}, \text{mix}\}$: $V_{\text{WIS}} = -1.339$ vs. default -1.383 , paired-bootstrap $\Pr(\text{WIS} \geq \text{default}) = 0.99$, recovering 28% of the gap-to-oracle. *Recipe: exclude deterministic-argmax policies from any WIS deployment set.*

APPENDIX F FQE BIAS AND SLOPE-CORRECTED IMPOSSIBILITY




Fig. 9. Fitted-Q estimator shrinks toward zero on the $N=9$ LOSO OPE corpus. Linear fit $V_{FQE} = 0.57 V_{GT} + 0.70$ ($r = +0.62$); weak-decodability subjects (red) exhibit 44% larger shrinkage.

FQE shrinks toward zero ($V_{FQE} = 0.57 V_{GT} + 0.70$, $r = +0.62$; 44% more on weak-decodability subjects, see Fig. 9). Any monotone affine correction $\hat{V} = (V_{FQE} - b)/m$ preserves within-subject argmax over policies, so the LOSO-corrected Δ versus naive FQE on policy selection is identically zero ($+0.000$) on all 9 subjects. Subject-conditioning is therefore *necessary* for any OPE-bias correction to alter cross-subject policy selection.

APPENDIX G HYPERPARAMETERS

TABLE XII
HYPERPARAMETERS ACROSS ALL EXPERIMENTS UNLESS NOTED.

Parameter	Value
Window length L / stride Δ	1.0 s / 0.25 s
Reward $R^+ / R^- / R_a / R_{\text{recal}}$	+1 / -5 / -0.5 / -0.3
Time cost λ	0.05 s^{-1}
EEGNet F_1 , D / params	8, 2 / $\sim 1.5\text{K}$
EEGNet pretrain epochs / lr / wd	20 / 10^{-3} / 10^{-4}
LaBraM params / fine-tune epochs / lr	5.8M / 8 / 10^{-4}
EEG-Conformer emb / depth / heads / drop	40 / 6 / 10 / 0.5
EEG-Conformer epochs / lr / β_1	80 / 2×10^{-4} / 0.5
Q-net hidden / depth	256 / 3
CQL α / CVaR α_{CVaR}	1.0 / 0.25
CMDP ϵ / dual-lr	$0.10 / 5 \times 10^{-3}$
γ / Polyak τ	$0.99 / 5 \times 10^{-3}$
Gradient steps (within / LOSO)	20K / 10K
Batch size	256
FQE Bellman iters / Bootstrap B	500 / 10,000

APPENDIX H REPRODUCIBILITY

Hardware: NVIDIA RTX 3080 (10 GB), CUDA 12.4; ~ 24 GPU-hours total. **Software:** Python 3.11, PyTorch 2.6, MNE, MOABB, braindecode 0.9, d3rlpy, scikit-learn. **Datasets** (public, via MOABB): BCI-IV-2b (BNCI2014_004), BCI-IV-2a (BNCI2014_001), Lee2019 (Lee2019_MI). All $\pm$ values are over 5 seeds; data preprocessing is deterministic. **Code** (anonymised, de-anonymised under MIT at camera-ready): pipeline + per-experiment summary JSONs at <https://anonymous.4open.science/r/NeuroPolicy-AFB8>; every claim traces to a JSON in experiments/.